NANYANG TECHNOLOGICAL UNIVERSITY SINGAPORE

# Coaching - ECS

Cloud Infrastructure Engineering

**Nanyang Technological University & Skills Union - 2022/2023**

# ECS Overview

Amazon Elastic Container Service (Amazon ECS) is a fully managed container orchestration service that simplifies your deployment, management, and scaling of containerized applications.

# Use Cases

## Modernize applications

Empower developers to build and deploy applications with enhanced security features in a fast, standardized, compliant, and cost-efficient manner with Amazon ECS.

## Automatically scale web applications

Automatically scale and run web applications in multiple Availability Zones with the performance, scale, reliability, and availability of AWS.

## Support batch processing

Plan, schedule, and run batch computing workloads across AWS services, including Amazon Elastic Compute Cloud (EC2), AWS Fargate, and Amazon EC2 Spot Instances.

## Train NLP and AI/ML models

Train natural language processing (NLP) and other artificial intelligence (AI) / machine learning (ML) models without managing the infrastructure by using Amazon ECS with AWS Fargate.

# Amazon Elastic Container Service Layers

**Provisioning**

Amazon Web Services Command Line Interface | Copilot | Management console | Amazon Web Services Cloud Developer Kit | Amazon Web Services Software Developer Kit

**Controller**

Amazon ECS scheduler

**Capacity options**

Amazon EC2 instances | Amazon Web Services Fargate | On-premises compute

# Amazon ECS Application Lifecycle

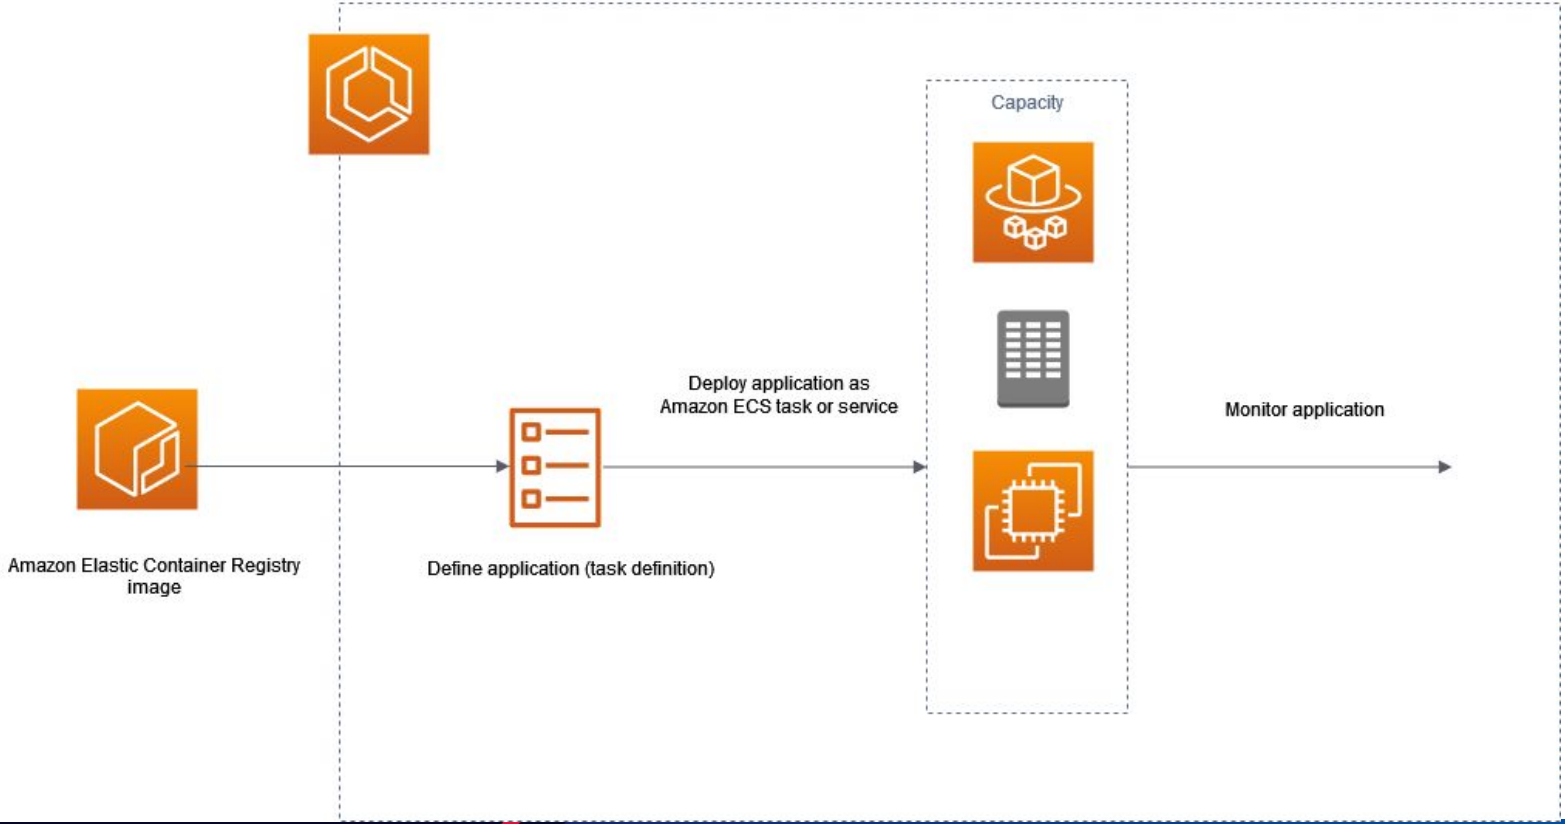

# ECS Terminologies

# ECS Cluster

An Amazon ECS cluster is a logical grouping of tasks or services. In addition to tasks and services, a cluster consists of the following resources:

- The infrastructure capacity which can be any of the following:
  - Amazon EC2 instances in the AWS cloud
  - Serverless (AWS Fargate (Fargate)) in the AWS cloud
  - On-premises virtual machines (VM) or servers (Using **ECS Anywhere:** Self-Managed Containers - Amazon ECS Anywhere - AWS)

# ECS Capacity Provider

**Fargate**: With Amazon ECS on AWS Fargate capacity providers, you can use both Fargate and Fargate Spot capacity with your Amazon ECS tasks. With Fargate Spot, you can run interruption tolerant Amazon ECS tasks at a rate that's discounted compared to the Fargate price. Fargate Spot runs tasks on spare compute capacity. When AWS needs the capacity back, your tasks are interrupted with a two-minute warning.

**EC2**: With AWS ECS on EC2 capacity providers, you can specify the auto scaling group and the setting required for the autoscaling group.

# AWS ECS Task Definition

A *task definition* is a blueprint for your application. It is a text file in JSON format that describes the parameters and one or more containers that form your application.

The following are some of the parameters that you can specify in a task definition:

- The Docker image to use with each container in your task
- How much CPU and memory to use with each task or each container within a task
- The launch type to use, which determines the infrastructure that your tasks are hosted on
- The Docker networking mode to use for the containers in your task
- The logging configuration to use for your tasks
- Whether the task continues to run if the container finishes or fails
- The command that the container runs when it's started
- Any data volumes that are used with the containers in the task
- The IAM role that your tasks use

# ECS Service vs Tasks

**Service** : Typically used for long-running tasks / containers that can be stopped and restarted. E.g. A web application

**Task**: A standalone task that is used to run a short running task. For e.g. A batch job, or a scheduled cron job to run at a specific time everyday and then exit

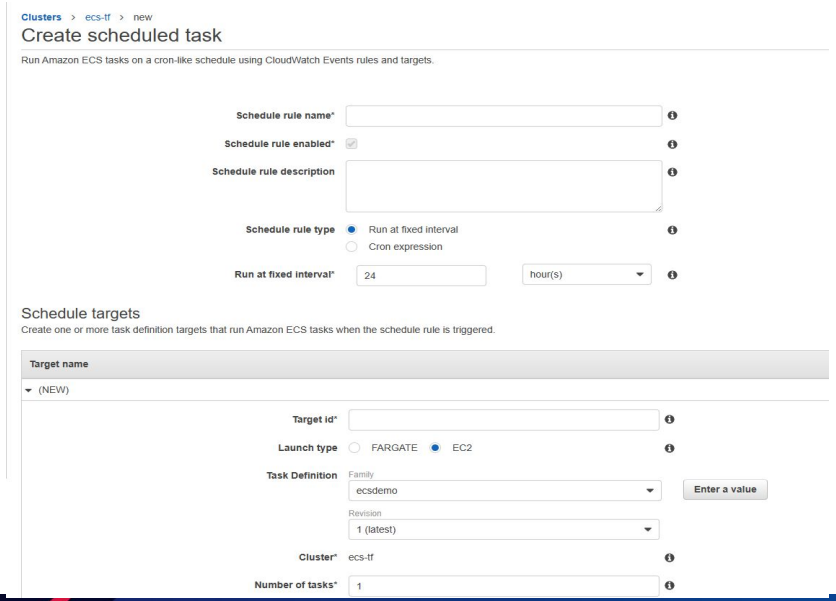

# ECS TaskExecutionRole vs TaskRole

**TaskExecutionRole**: IAM role that executes ECS actions such as pulling the image and storing the application logs in cloudwatch.

**TaskRole**: IAM role used by the task itself. For example, if your container wants to call other AWS services like S3, SQS, etc then those permissions would need to be covered by the TaskRole

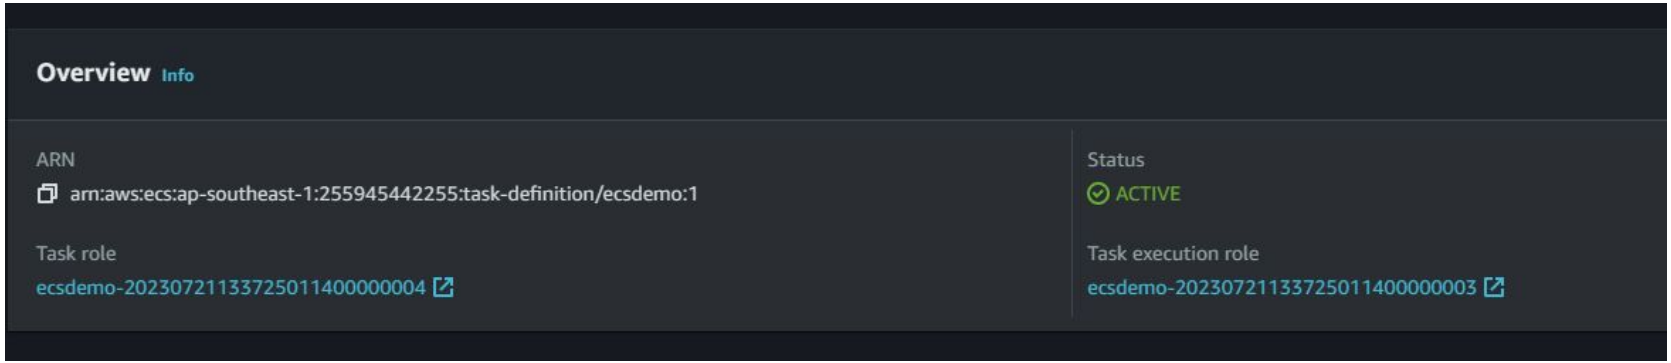

# Activity 1: Create ECS cluster and task definition using a default image

- Create a ECS cluster with Fargate type
- Create a task definition with the following specs (leave all else as default other than the ones mentioned below):
  - Name: yourname-httpd
  - Under container-1
    - Name: httpd
    - image uri: `public.ecr.aws/docker/library/httpd:latest`

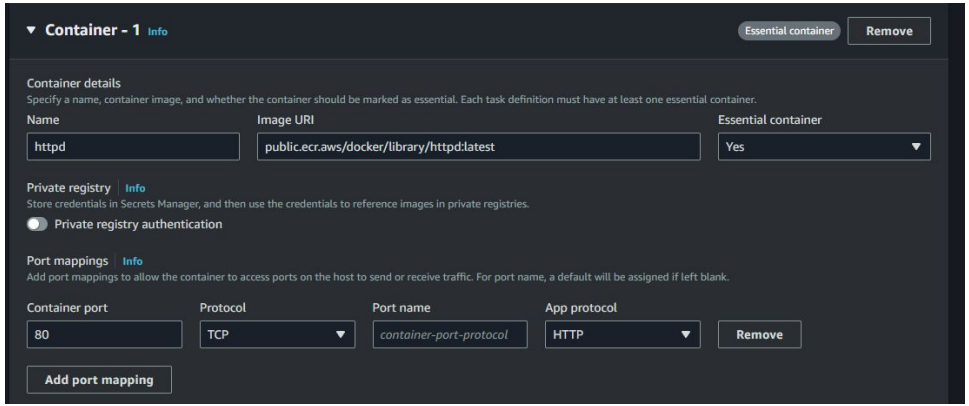

# Task Definition JSON

- You can view this after you've finished creating your task definition. (Task definitions -> Click on yours -> JSON)

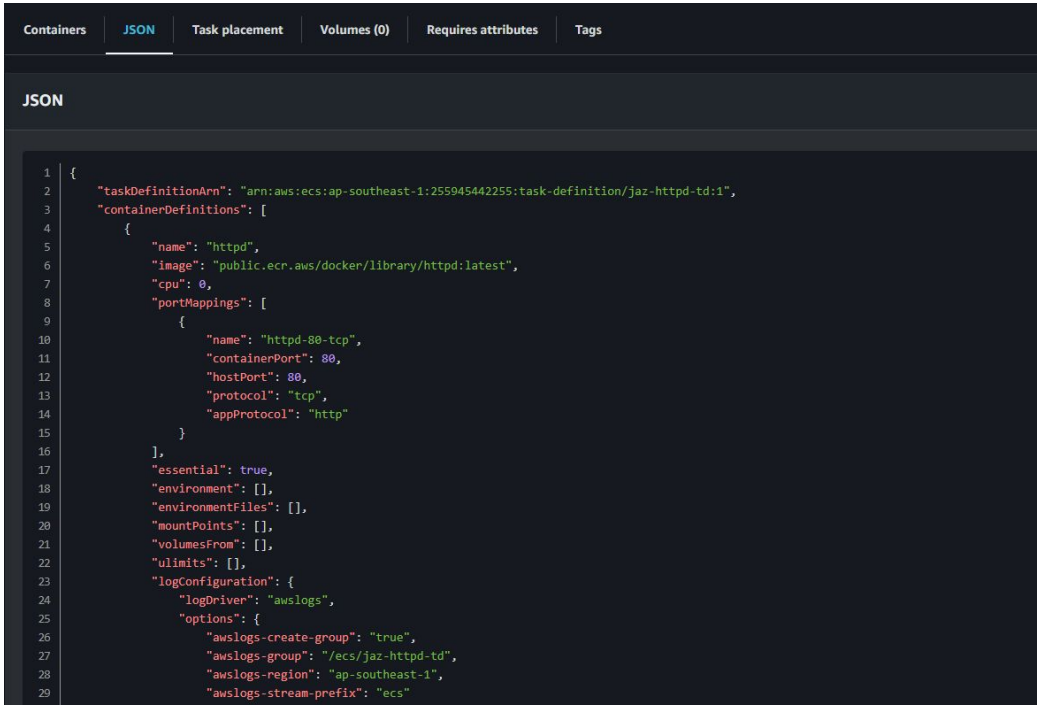

# Activity 1: Create ECS service

- Go to the ECS cluster that you've created and go to the **Services** tab.
- Click on **Create (Can refer to next slide)**
- Under Deployment Configuration
  - Family: Specify the task definition which you created in the previous slide
  - Service name: httpd
- Under Networking
  - Create your own security group with inbound port 80
  - For VPC: Feel free to use any VPC, but remember to use a public subnet
  - Public IP: Turned on

# Service Creation

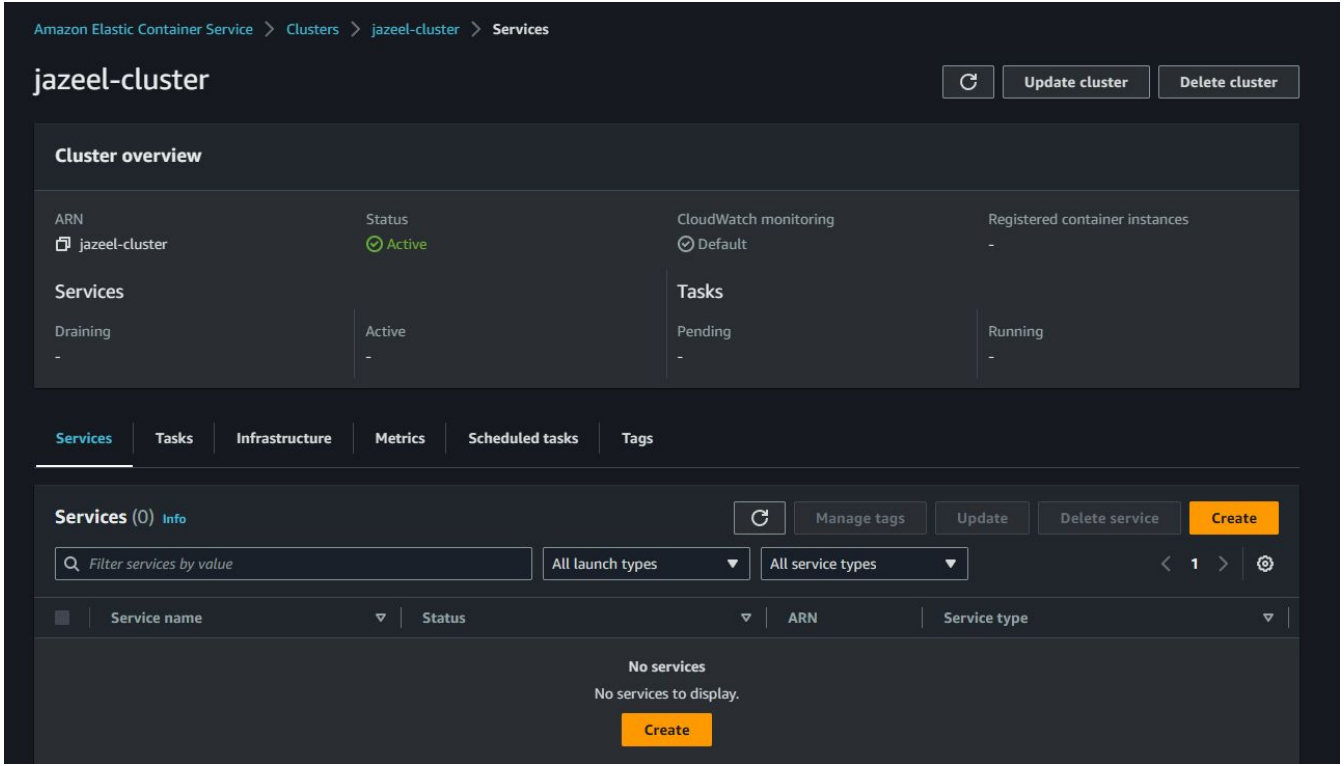

# Activity 1: Create ECS cluster and service using a default image

- Service configurations:

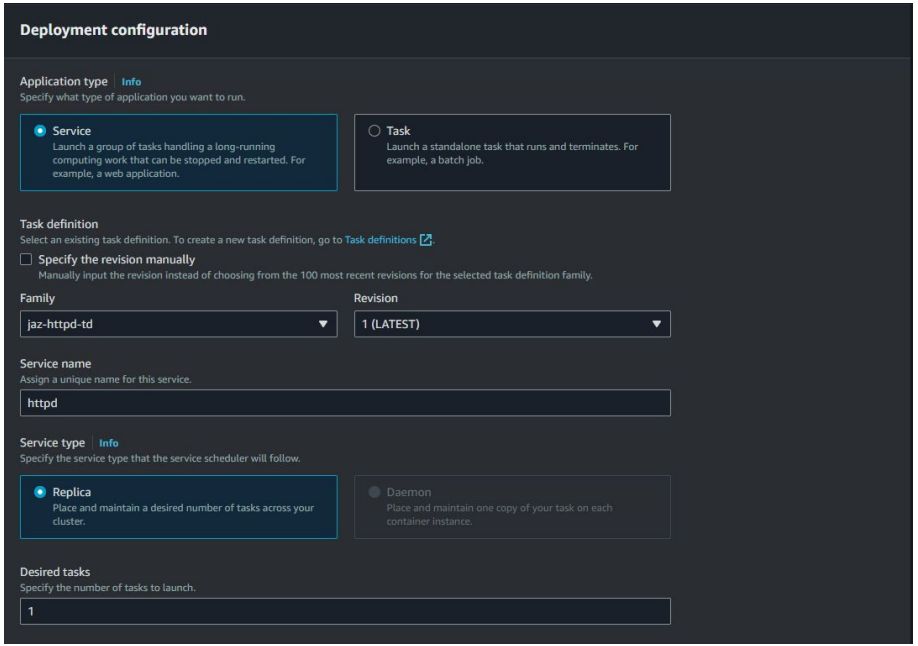


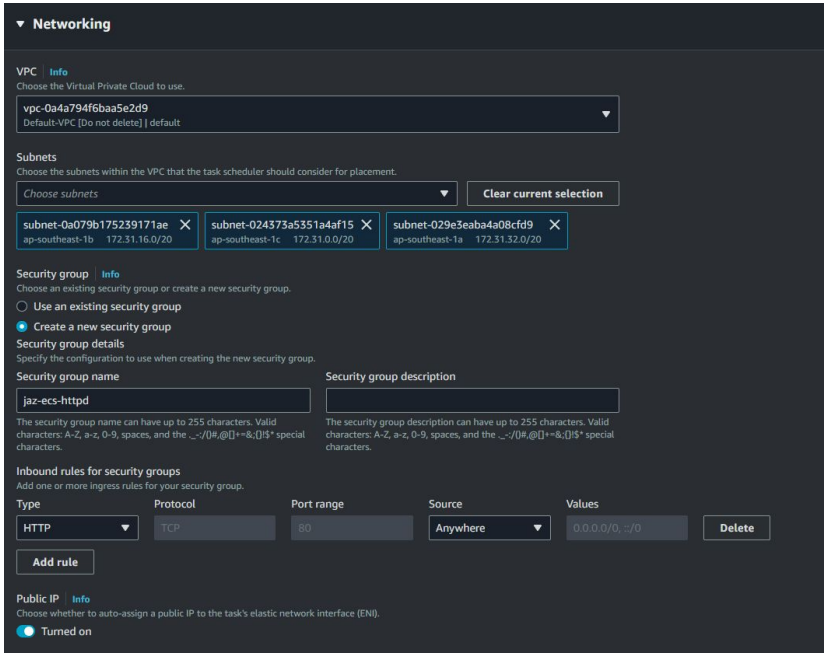

# Activity

1) Access the cluster you have created
2) Go to the **tasks** tab
3) Click on the task
4) Obtain the public IP and try to access it in a browser tab
5) Delete the service that you've created.
6) Retain the ECS cluster for the next activity

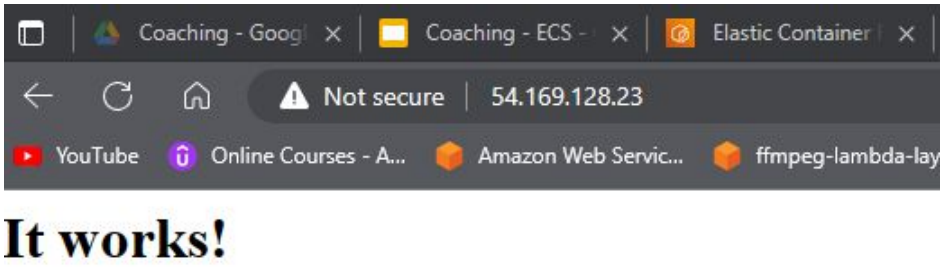

# Service - Health and metrics

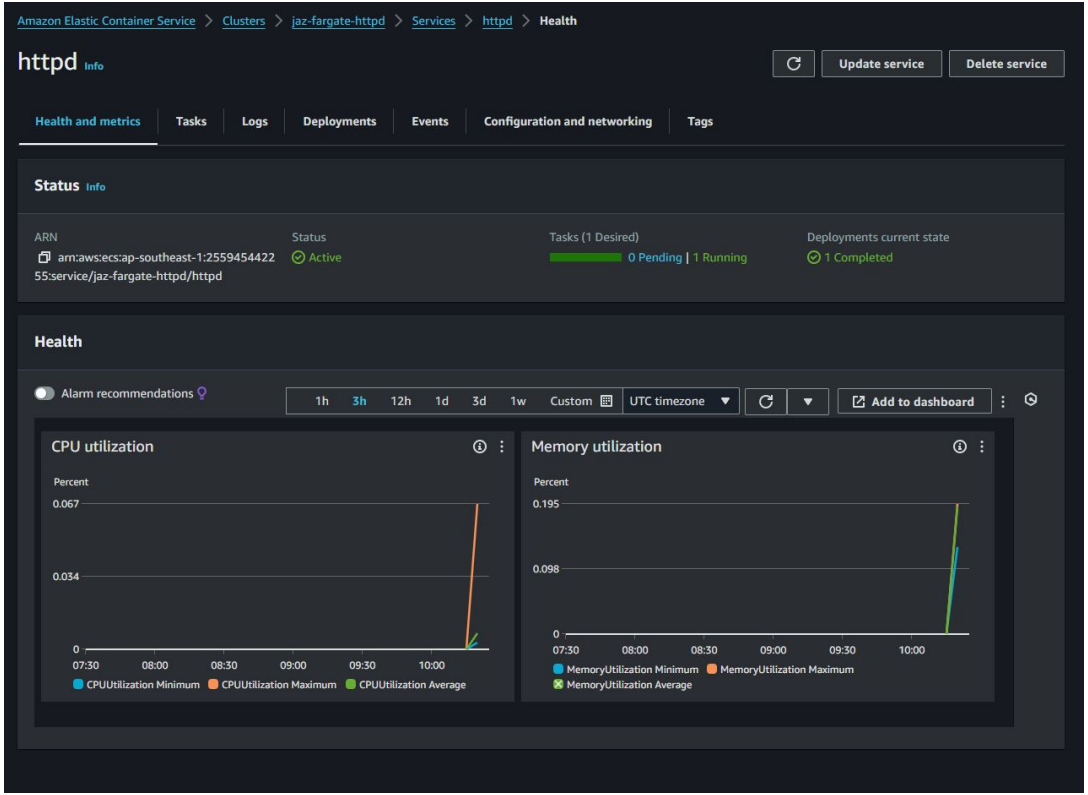

# Service - Logs

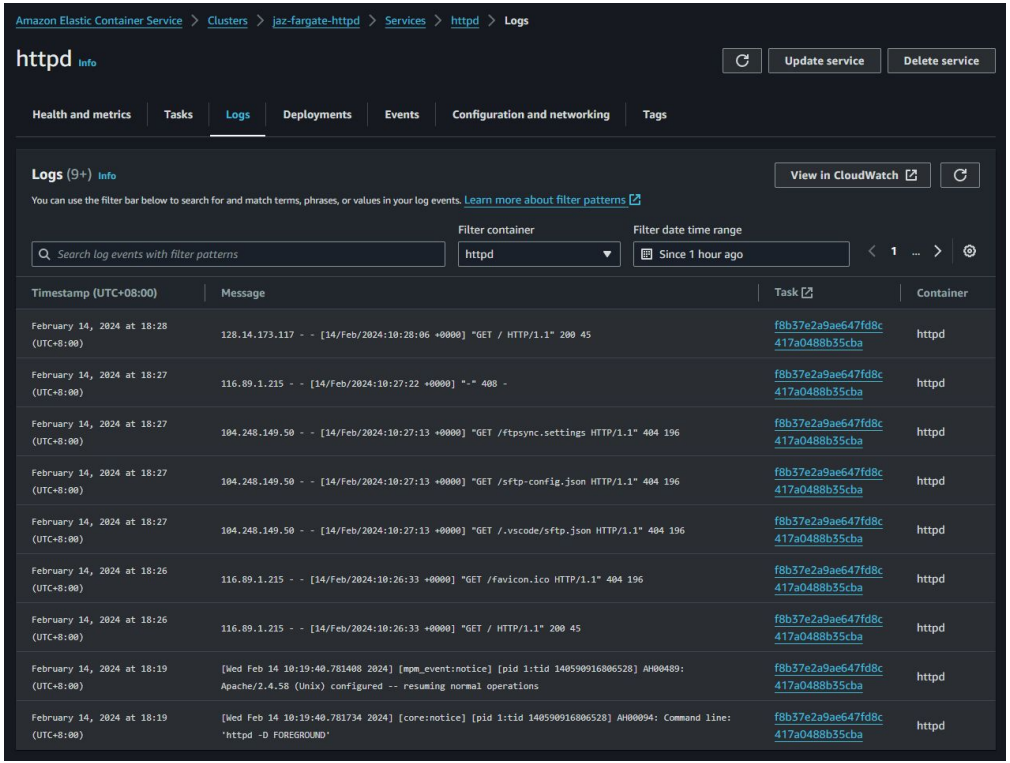

# Service - Deployment

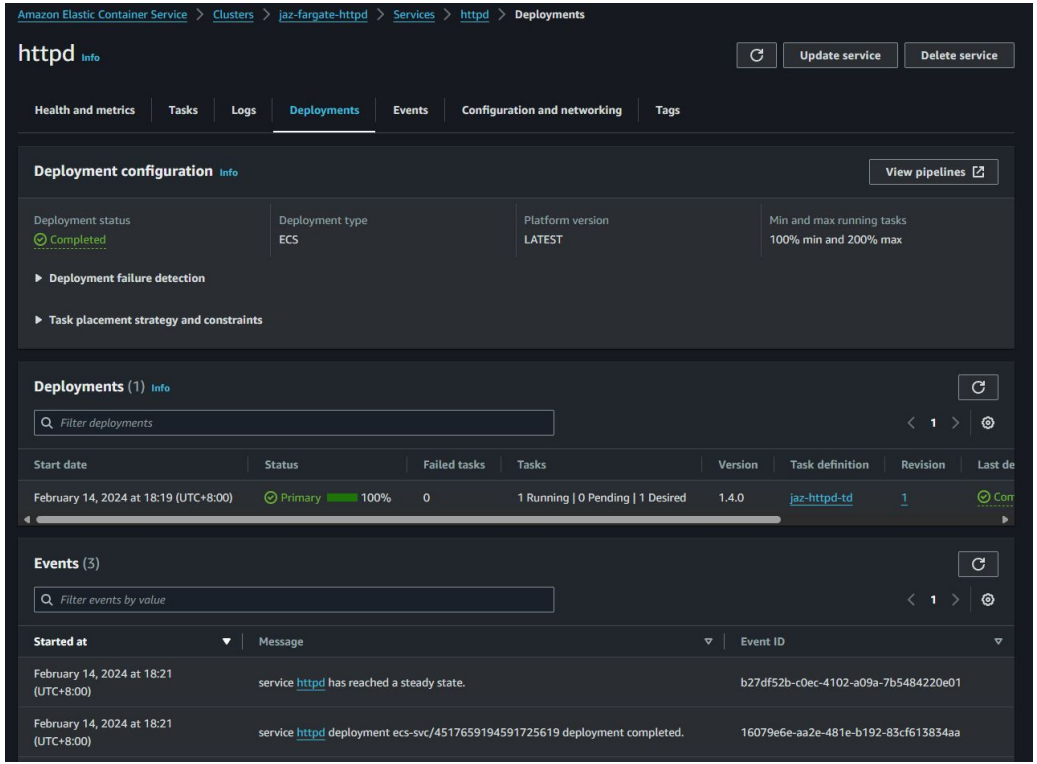

# Service - Networking

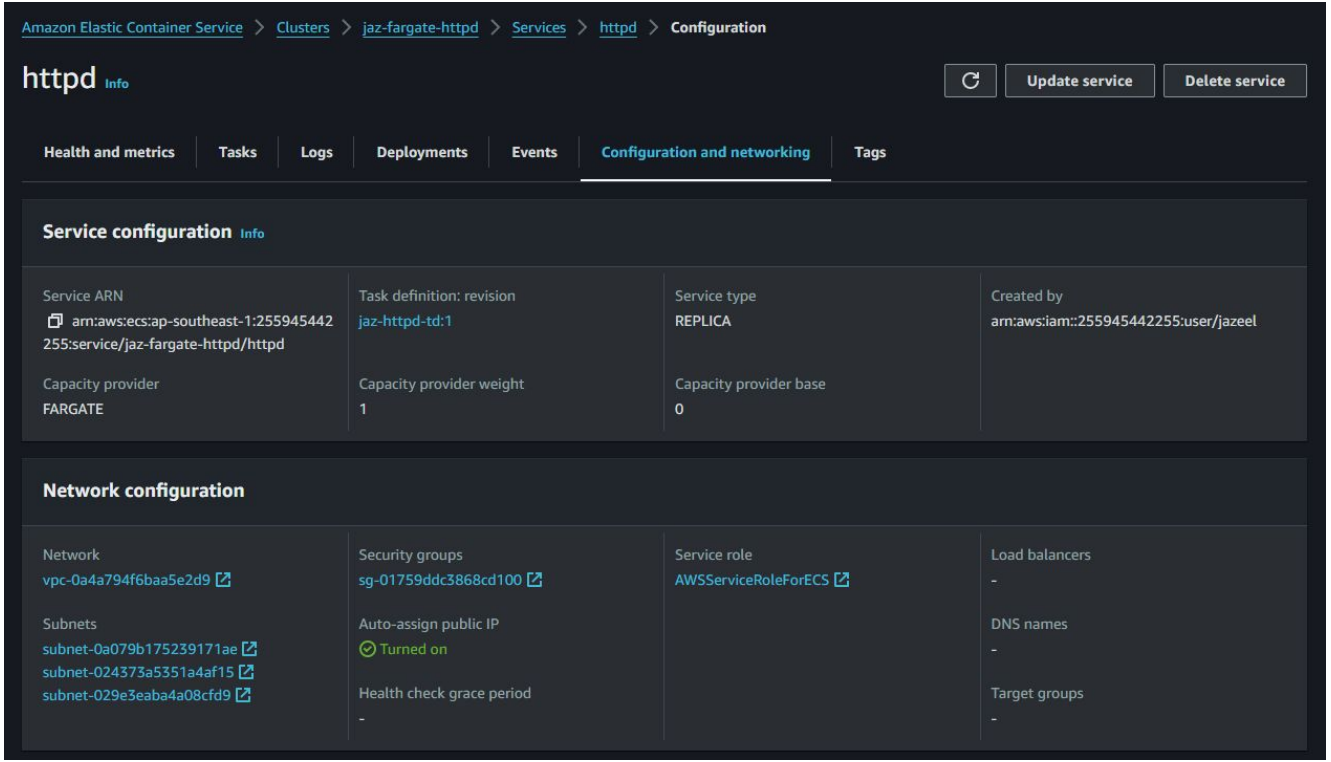

# Post-Activity for Activity 1

1) Delete the service that you've created.
2) Retain the ECS cluster for the next activity

# Activity 2: Using custom image in ECS

- Create a new private repository in AWS ECR
- Using the sample code in the next page, build your own docker image and push it to that ECR (Feel free to also try with other docker images or code if you have)
- Create a new ECS task definition using the custom image you've pushed to your ECR (You can obtain the image URI from your repo). Remember to also update the container port to be the same port as your docker image
- Create a new service in your ECS cluster with this new task definition, with the respective security group rules required.

# Activity 2: Using custom image in ECS

- app.py - hello-flask/app.py at main · jaezeu/hello-flask · GitHub
- requirements.txt - hello-flask/requirements.txt at main · jaezeu/hello-flask · GitHub
- Dockerfile - hello-flask/Dockerfile at main · jaezeu/hello-flask · GitHub

If you do not already have, Create your own github repository with the files above. Files you should have in your git repo:

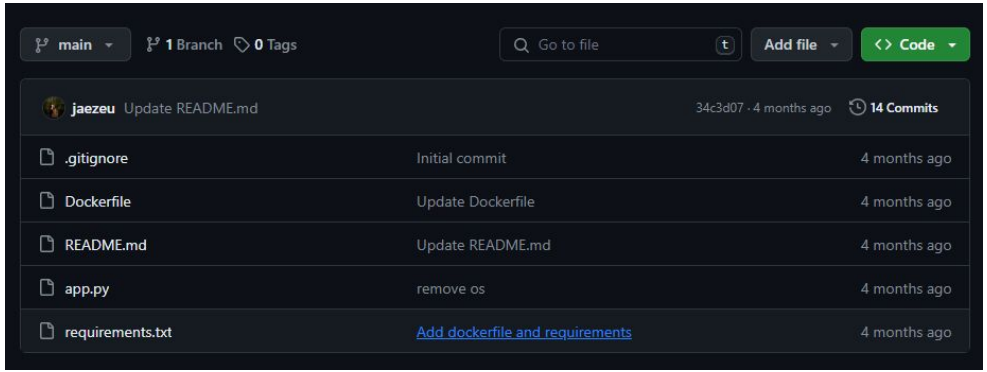

# Post-Activity for Activity 2

1) Delete all the resources you've created in today's activity:
   a) Security groups
   b) ECS cluster + Task definition + Service
   c) ECR repository

# Post-Activity

Some resources you can refer to (which would help you write lesser code), in order to create all of our today's infra:

ECS cluster + Task definition + Service: terraform-aws-modules/terraform-aws-ecs: Terraform module to create AWS ECS resources 🇺🇦 (github.com)

ECR Repository: aws_ecr_repository | Resources | hashicorp/aws | Terraform | Terraform Registry

ECS security group: terraform-aws-modules/terraform-aws-security-group: Terraform module to create AWS Security Group resources 🇺🇦 (github.com)